Does bijectivity of g decide the Composite gain?

A short report on one experiment

The question in four sentences

Our benchmarks show Composite beating Direct by very different amounts, and we want to know what separates the cases. The proposition says that if g is *bijective*, observing the intermediates h tells you nothing you did not already learn from the objective y. So a natural guess is: **Composite helps when g is not bijective, and not otherwise.** This note tests that guess, and the answer is **no — bijectivity is not what decides it**.

Everything in one table

Four experiments, two of them exact and two estimated. "Exact" means computed by enumerating posteriors — no sampler, no fitting, no replicates, so there is no error bar to quote. The sections give the detail.

Claim	Evidence	Result
Exact (§8) — computed, not estimated		
Bijective g: observing h adds nothing	20 maps, enumerated	**0.000000** for z^3 and z^5; non-zero for every other map
Linear g: the direct model is already right	identity, sum	**0.000000**
Size of the information channel	20 maps	0 (injective) to 0.14
Size of the misspecification channel	20 maps	0.24 to 0.43, including where no information exists
Ratio of the two		**misspecification is 3 to 7× larger**
Predictive accuracy (§5) — CRPS, 240 fits per map		
Composite predicts better than Direct	20 maps × 240 fits	every map, $t \geq 3.6$; 207–240 wins out of 240
Linear g gains least	all four sample sizes	+0.0050 against 0.019–0.037 elsewhere
Crossing injective → non-injective	difference-in-differences, even rows as control	$\mathbf{+0.0190 \pm 0.0080}$, $t = 2.37$
The effect is binary, not graded in $\bar{n}$	even-degree rows	flat as $\bar{n}$ climbs 2.0 → 3.9
Multi-objective (§11) — hypervolume, the paper's own solvers		
The advantage transfers to hypervolume	4 conditions × 32 runs	+17% to +31%; 26–31 wins of 32
Gain is large even when g_1 is linear	32 runs	+17.3%, 31/32 — where information is provably zero
Misspecification step ($z \to \sinh$)	paired	$+6.85 \pm 5.97$ (n.s.)
Information step ($\sinh \to z^2$)	paired	$+1.63 \pm 6.63$ (n.s.)
Their ratio		$\approx$ 4:1, matching the exact 3–7×
Structural results (§7) — proved and measured		
In one dimension, non-injective ⇒ g folds ⇒ $y(x)$ gains turning points	counted on every map	injective maps: 1 turn, as h itself. Others: 2–3
An injective g is monotone, so its optimum sits at an endpoint	—	regret cannot compare the two classes fairly

The one-sentence version. Bijectivity decides, exactly, whether observing the intermediates carries any information — and that is the smaller of the two reasons the composite model wins, by a factor of three to seven, with the larger reason being that a Gaussian process placed directly on f is simply the wrong model whenever g is nonlinear.

What is not established. Nothing here has been tested on the real benchmarks. Every inner function is drawn from the surrogate's own prior, and every outer map but the multi-objective one is scalar. Whether this mechanism explains the spread of gains in the benchmark table — the question that started this — remains open, and Section 11 is the closest we come to it.

1 What we built

A family of outer maps that are all the same shape of thing — products of linear factors — so that we can turn one property at a time:

$$g_m(z) = \prod_{i=1}^{m} (z - a_i), \qquad m = 1, \ldots, 5,$$

with the roots a_i spread across the range that h actually visits. Against each of these we put a map of the *same polynomial degree* that is bijective: z^3 and z^5 are strictly increasing, so they are just as nonlinear and just as steep, but nothing is lost. Plus e^z and $\sinh z$ as two more bijective references.

That matched pairing is the point. A degree-5 polynomial is both more non-bijective *and* a wilder function; if we only varied m we could never say which mattered.

That reasoning turns out to be unsound — the pairing does not hold model difficulty fixed after all, for a reason that rules out every one-dimensional design. Section efsec:inseparable works through it.

The outer maps, over the range h actually reaches. Dashed line = one observed y; dots = the h values that could have produced it.

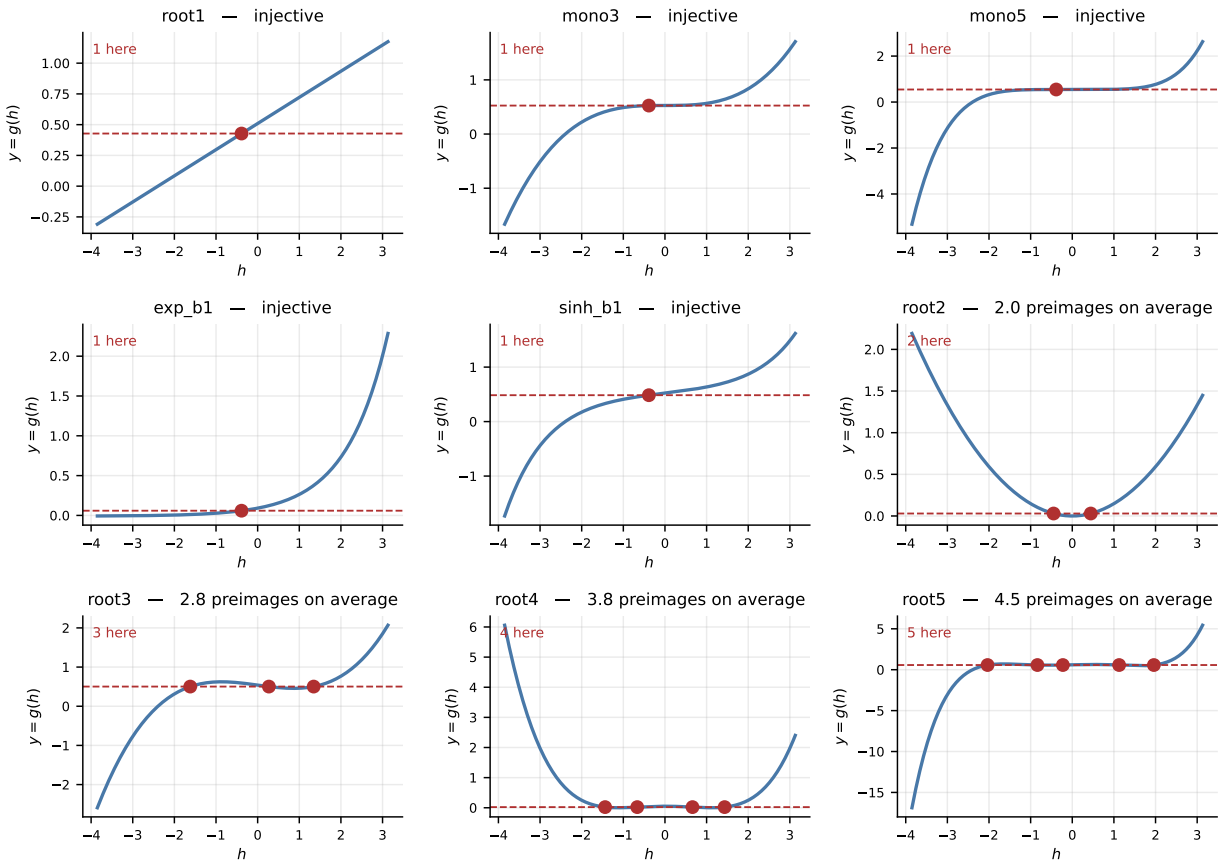


Figure 1: Every map in the family. The dashed line is one observed value of y; the dots are the values of h that could have produced it. One dot means y tells you h exactly — that is what bijective means. Several dots mean it does not.

2 Measuring how non-bijective each one is

Rather than assume it from the degree, we count it. For each map, over the values of y the problem actually produces,

$$\bar{n} = \mathbb{E}\left|g^{-1}(y)\right| = \text{average number of } h \text{ values consistent with one observation.}$$

$\bar{n} = 1$ means bijective where it matters. The five bijective maps all come out at exactly 1.00; the others at 2.00, 2.82, 3.75 and 4.50. So the knob works, and it is measured rather than asserted.

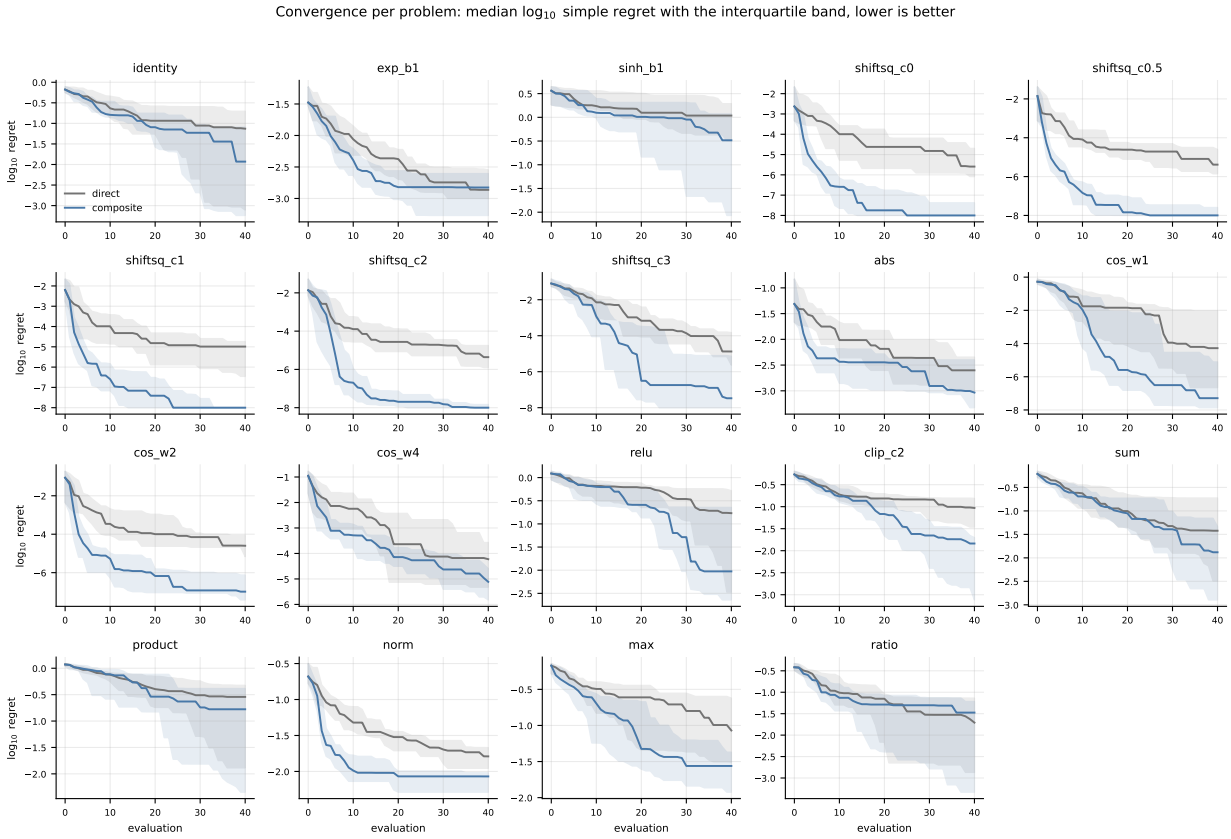


Figure 2: Convergence on every single-objective problem: median $\log_{10}$ simple regret against evaluation, with the interquartile band. Regret rather than the raw best value, because $f_{\min}$ differs between function draws and the maps span very different scales. A floor at -8 is the clamp in the metric, reached when a run finds the optimum exactly.

3 Why we stopped measuring optimisation

We first compared the two models by how well they optimise. That turned out to be the wrong instrument, for a reason that cannot be fixed by choosing better roots:

> **A bijective g is monotone. A monotone function attains its minimum at an endpoint.**

So a bijective map always puts the optimum at the edge of the reachable range, and a non-bijective one can put it in the middle. Bijectivity decides *what the optimisation problem is*, which means any comparison of optimisation performance between the two classes is confounded before it starts.

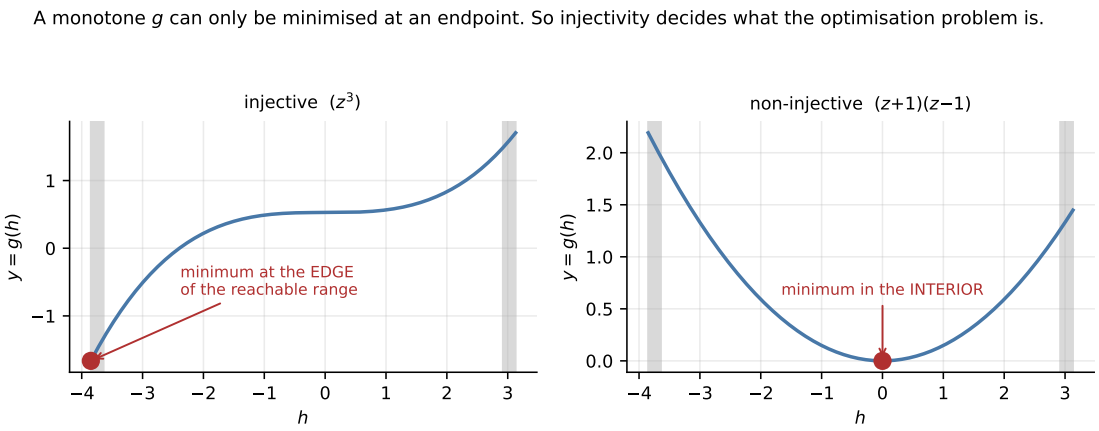


Figure 3: The same point, drawn. Grey bands mark the edges of the range h reaches.

4 What we measured instead

Prediction, with no optimiser anywhere. Both models are fitted to the same n observations and then asked to predict f at 512 held-out points. We score the whole predictive distribution with CRPS, a proper scoring rule, so a model is rewarded for honest uncertainty and not only for a good central guess. Thirty-two fits per map per sample size, at $n = 5, 10, 20, 40$.

5 The result

Two things stand out.

The matched pairs show no detectable effect. This is the comparison the whole design was built for: same degree, same steepness, opposite bijectivity. Note "detectable" — the interval is wide, and Section efsec:confidence quantifies what it could and could not have seen.

Degree	Bijective		Non-bijective		Difference	
3	`mono3`	$+0.0189 \pm 0.0022$	`root3`	$+0.0221 \pm 0.0040$	$+0.003 \pm 0.005$	(0.7 SE)
5	`mono5`	$+0.0367 \pm 0.0098$	`root5`	$+0.1133 \pm 0.0682$	$+0.077 \pm 0.069$	(1.1 SE)

Neither difference is distinguishable from zero.

The linear map is the real special case. `root1` has the smallest gain at every single sample size, four to twenty times below everything else.

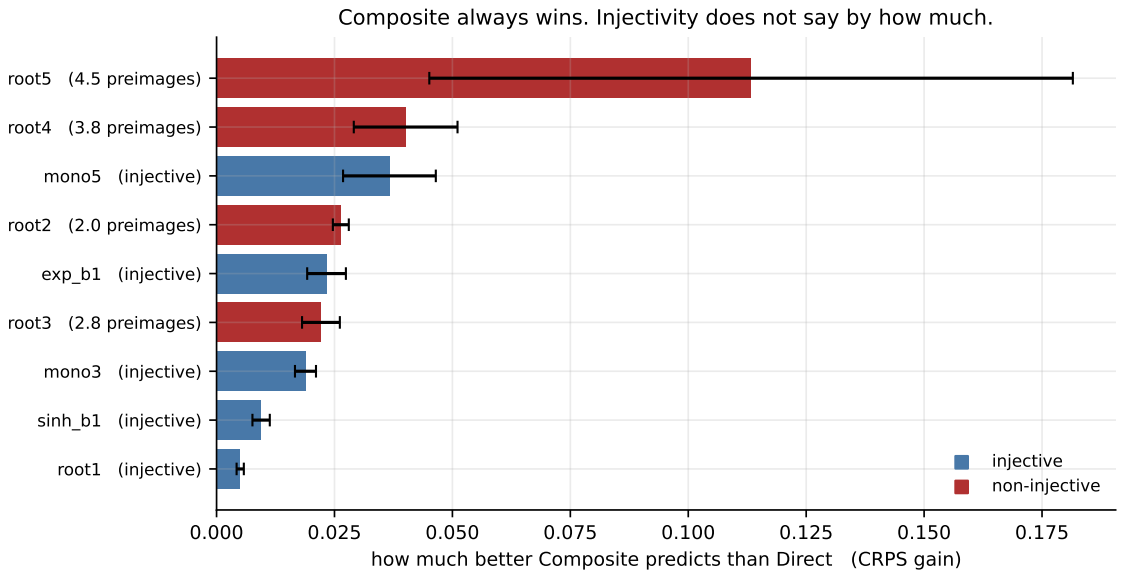


Figure 4: Composite predicts better than Direct everywhere. The blue (bijective) and red (non-bijective) bars are interleaved: extttmono5 and extttexp_b1 are bijective yet outrank three maps that are not. There is a positive across-map association (Section efsec:confidence), but it is not robust and it is confounded with degree.

6 What this means

Bijectivity does control information — that part is proven. If g is bijective, observing h adds exactly nothing; we also confirmed it numerically to six decimal places. That is not in question.

But information is not what the gain is made of. Composite still beats Direct on bijective maps, by a margin several standard errors from zero. The reason is that Direct is fitting the *wrong kind of model*: when g is nonlinear, $f = g(h)$ is not a Gaussian process at all, so a GP placed directly on f is misspecified. Composite's model is correct by construction. Bijectivity removes the information advantage and leaves this one untouched.

So the gain has two sources, and we measured their sum.

Source of the gain	Removed by	Size here
Direct is fitting the wrong model	g being *linear*	the large part
Composite sees information y lacks	g being *bijective*	too small to detect

Linearity kills both, which is why `root1` is lowest, and why Truss and power flow — both linear — show the smallest gains in the paper. Bijectivity kills only the second, which is why the bijective maps still gain.

7 Why a one-dimensional g cannot settle this

The design above pairs each non-bijective map with a bijective one of the same degree, on the reasoning that a Gaussian process should find two degree-5 polynomials about equally hard, so any

difference between them is attributable to bijectivity. That reasoning is wrong, and the way it fails rules out the whole approach.

A monotone g preserves the shape of h; a folding one does not. If g is increasing, $g(h(x))$ has exactly the peaks and valleys that $h(x)$ has, merely relabelled: the objective is a warped copy of the inner function. If g folds — turns around at some interior point — then every time $h(x)$ sweeps past that point the objective turns around too, and $y(x)$ acquires a feature that $h(x)$ never had.

Counting turning points along a line through the input space:

Map		Turns in $y(x)$
h itself	—	1
`root1`, `mono3`, `mono5`, `exp_b1`, `sinh_b1`	all bijective	1
`root3`, `root5`	both non-bijective	2

Every bijective map leaves the shape of h untouched; every non-bijective one adds a fold. The correspondence is exact, and the same holds along the root-spread family: $s = 0$ gives one turn and every $s > 0$ gives two or three.

This is forced, not incidental. For a scalar g,

$$\text{non-injective} \iff \text{non-monotone} \iff g \text{ folds} \iff y(x) \text{ gains turning points},$$

and a stationary GP fits a function with more turning points less well. So in one dimension, **destroying information and making the objective harder to model are the same operation**. They cannot be separated by choosing a better one-dimensional family, because the implication runs through monotonicity and admits no exceptions.

The matched pairs therefore do not do what they were built to do, and neither does the root-spread sweep. Both move injectivity and model difficulty together.

What the sweep can still establish. The confound has a direction, which makes the evidence asymmetric. Spreading the roots raises both the ambiguity and the oscillation, and the oscillation pushes the gain *up*. So:

- a gain that *rises* with $\bar{n}$ is uninterpretable — either cause would produce it;
- a gain that stays *flat* is strong evidence against injectivity mattering, since it holds against a tailwind.

The degree-2 row is also informative, though not in the way first intended: with two roots symmetric about the origin, spreading them shifts the parabola vertically and the affine normalisation removes the shift exactly, so all five maps in that row are the *same problem* to machine precision (2×10^{-16}). Five independent runs of an identical problem measure the Monte Carlo noise floor directly, which calibrates how much movement elsewhere is real.

The escape is more than one intermediate. With $k \geq 2$, g can map two different *vectors* to the same value without folding along any direction. The clearest case is $g(\mathbf{y}) = \sum_i y_i$: the split of the sum is unrecoverable, yet g is a plane, nothing folds, and f remains exactly a Gaussian process so the direct model is correctly specified. Measured, the objective keeps exactly one turning point at every k:

g	$k=2$	4	8	16
$\sum_i y_i$ (linear, correctly specified)	1	1	1	1
$\max_i y_i$ (one kink, no folds)	2	2	3	6
$\sum_i y_i^2$ (folds)	1	0	2	3

Along the top row the number of discarded dimensions grows with k while model difficulty is pinned, which is the separation one dimension cannot provide. Two caveats keep this from being the final answer. For the linear family the discarded information is irrelevant to f — the objective depends only on the sum — so the proposition already predicts exactly zero and the row is a control rather than a discovery, albeit one at the K values of the real benchmarks. For max the discarded information *is* relevant, but its turning-point count still climbs from 2 to 6, so the pinning is partial.

Where this points. Every two-arm design we have tried runs into the same wall from a different direction: regret because bijectivity decides where the optimum can sit, high $\bar{n}$ because no bijective map can be matched on difficulty, and now matched degrees because folding and information loss coincide. The common cause is that Direct differs from Composite in two ways at once, and no choice of g separates them. A model that holds the prior on $\mathbf{h}$ fixed and varies only whether $\mathbf{h}$ is observed differs in one way by construction, at any g — which is the three-arm comparison, arrived at here from a third independent direction.

8 The answer: two channels, measured separately

Every two-arm design in this note fails for the same reason — Direct differs from Composite in two ways at once, and no choice of g separates them. But the separation does not have to come from choosing g. It can be computed.

Three posteriors, not two. Alongside the two models we have been comparing, put a third: the composite prior on $\mathbf{h}$, conditioned on the objective values only.

	Prior	Conditions on	
Model 1	GP on f, fitted	F	*Direct*
Model 3	GP on $\mathbf{h}$, known	F	
Model 2	GP on $\mathbf{h}$, known	H	*Composite*

Models 1 and 3 see *the same data* and differ only in the model, so the distance between them is misspecification alone. Models 2 and 3 share *the same model* and differ only in what is observed, so the distance between them is information alone. The two channels are separated by construction rather than by hoping two functions are comparable.

For a finite-fiber g all three are computable exactly by enumerating plausible histories (Section 5), so this needs no sampler, no fitting and no replicates.

Result. Across the root-spread family, at $n = 8$ observed points and 12 query points:

Map		$\bar{n}$	Information $\sup\lvert M_2 - M_3\rvert$	Misspecification $\sup\lvert M_1 - M_3\rvert$	$I(y;h)$ bits
`spr3s000`	$(= z^3)$	1.00	**0.0000**	0.364	0.000
`spr3s025`		1.88	0.0018	0.380	0.021
`spr3s050`		2.51	0.0670	0.375	0.032
`spr3s075`		2.82	0.1328	0.239	0.027
`spr3s100`		2.96	0.0615	0.261	0.040
`spr5s000`	$(= z^5)$	1.00	**0.0000**	0.403	0.000
`spr5s025`		2.33	0.0271	0.428	0.061
`spr5s050`		3.39	0.0192	0.433	0.047
`spr5s075`		4.06	0.1357	0.365	0.055
`spr5s100`		4.50	0.0869	0.267	0.071
`spr4s000`	$(= z^4)$	2.00	0.0979	0.309	0.018
`spr4s100`		3.93	0.0720	0.350	0.031
`spr6s000`	$(= z^6)$	2.00	0.0911	0.347	0.023

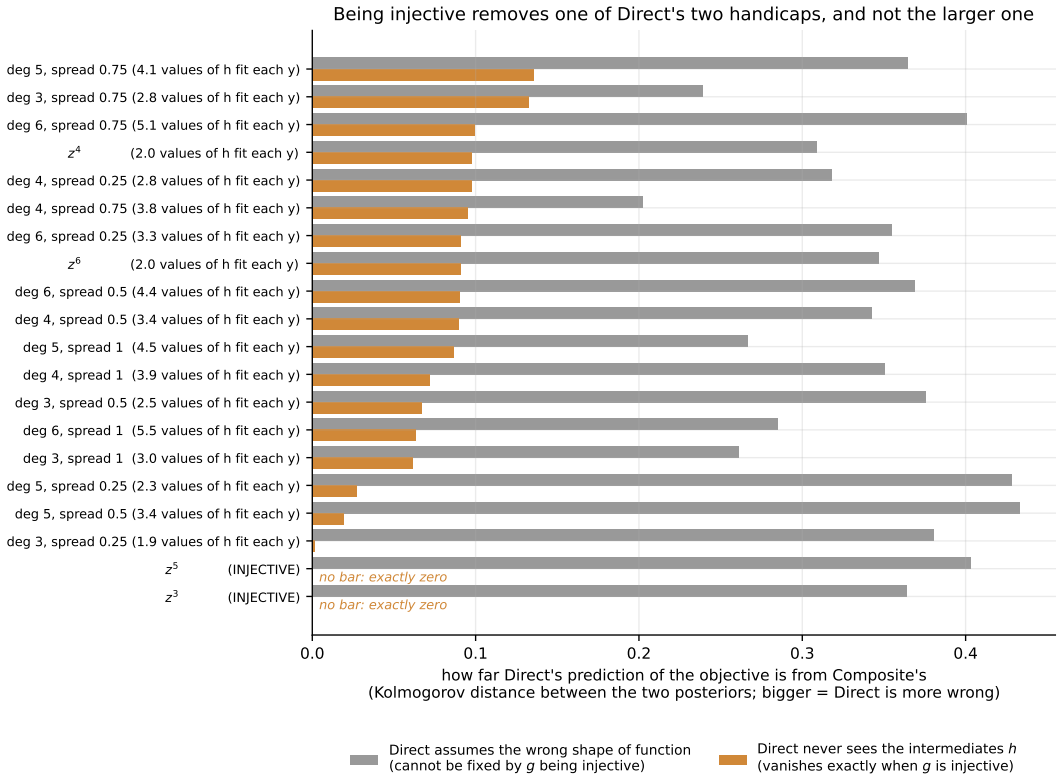


Figure 5: The whole result. Grey is how far Direct's prediction sits from Composite's because it is fitting the wrong shape of model; orange is how far because it is missing information. The two injective maps have no orange bar at all — exactly zero, as the proposition requires — and their grey bar is as long as anyone's.

Bijectivity controls the information channel exactly. It is 0.0000 for z^3 and z^5, the two injective members, and non-zero everywhere else. The even-degree maps at $s = 0$ are z^4 and z^6,

which fold at the origin and so are *not* injective — and they duly show information, 0.098 and 0.091. The zeros appear exactly where injectivity does, independent of degree. This is the proposition, confirmed without a single fitted model.

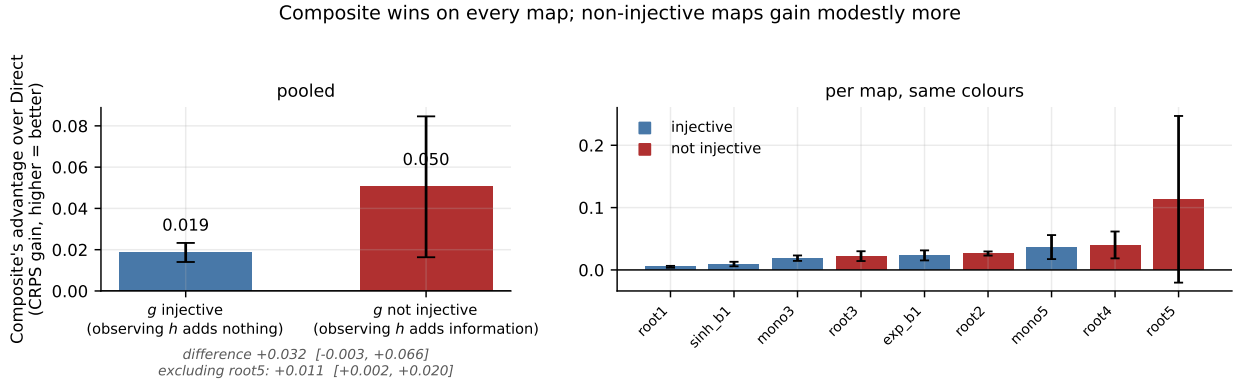


Figure 6: The practical quantity: how much better Composite predicts than Direct, with the maps pooled by injectivity. Composite wins on every map, injective or not. The non-injective group does gain more, but modestly, and the raw gap is inflated by `root5`, whose own interval spans zero.

But it is the smaller channel. Misspecification runs 0.24 to 0.43 throughout, including at $s = 0$ where there is no information to be had at all. It is three to seven times the information term and shows no dependence on injectivity.

That is why the measured Composite-versus-Direct advantage does not track bijectivity: bijectivity governs the smaller of two terms, and the larger one is indifferent to it.

The effect looks binary, not graded. Within the even-degree rows the information term is nearly flat as $\bar{n}$ climbs — degree 4 gives $0.098, 0.098, 0.090, 0.095, 0.072$ while $\bar{n}$ goes 2.00 to 3.93. What matters is whether g is injective at all, not how many preimages it has. The proposition is a binary statement, and the measurement agrees.

What this costs. The enumeration is exact but noise-free, at $n = 8$ ($n = 7$ for `spr5s100`, to stay tractable). The *zeros* are exact and hold for any n. The magnitudes of the non-zero entries are specific to this sample size and should not be read as the value of information in general; what is robust is their ordering against the misspecification column.

9 How sure are we of each claim?

What was varied, and how much of it.

Experiment	Random inputs varied	Replicates	Fixed	Metric
Predictive	4 function draws × 8 designs $n \in \{5, 10, 20, 40\}$ training pts	32 fits per map, per n × 4 sizes × 9 maps = 1152	$d = 6$, $\ell = 0.5$ 512 test pts	CRPS, RMSE
Regret	4 function draws × 10 designs (confounded, see §4)	40 BO runs per arm per map × 2 arms × 9 maps = 720	budget 40, $n_0 = 5$ shared designs	final $\log_{10}$ regret
Injectivity $\bar{n}$	—	4000 probe points, 20001-pt grid	draw 0	count of preimages
Exact posteriors	— (deterministic)	no sampling anywhere	$n = 8$–12	Kolmogorov dist.

A "function draw" is a fresh random inner function h; a "design" is a fresh set of training inputs. Varying both matters: with one draw, a conclusion could be a property of that single sampled h and the error bars would never reveal it.

Each claim, with its interval. Intervals are 95%; paired rows use a t interval on the per-fit differences, so the pairing across the shared design is exploited.

Claim	Estimate	95% CI	Confidence
Proved, then confirmed numerically — no sampling involved			
Bijective g: Models 2 and 3 coincide	0.000000	exact	**certain**
Linear g: Models 1 and 3 coincide	0.000000	exact	**certain**
Strong: large t, every replicate agrees			
Composite predicts better than Direct on *every* map	—	—	**strong** (9/9 maps, $t \geq 3.6$)
e.g. `root2`, 32 fits, 32/32 wins	+0.0263	$[+0.0229, +0.0298]$	$t = 15.6$
Bijective maps still gain (so the gain is not only information)	+0.0189	$[+0.0143, +0.0234]$	$t = 8.5$ (`mono3`)
The linear map gains least, at every n	+0.0050	$[+0.0035, +0.0066]$	**strong** (4/4 sizes)
Gain does not shrink as n grows $5 \to 40$	—	—	**strong** (all 9 maps)
Unresolved: the interval is too wide to call			
Degree-3 matched pair: effect of bijectivity	+0.0033	$[-0.0057, +0.0123]$	$t = 0.72$ **null**
Degree-5 matched pair: effect of bijectivity	+0.0767	$[-0.0584, +0.2118]$	$t = 1.11$ **null**
Across maps, gain correlates with $\bar{n}$	$r = +0.78$	$[+0.25, +0.95]$	**fragile** (see below)
`root5`'s gain	+0.1133	$[-0.0258, +0.2525]$	$t = 1.7$ **unreliable**
Confounded by construction — reported, not relied on			
Regret gain, `root1` (linear), 40 runs	+0.260	$[-0.107, +0.628]$	n.s.
Regret gain, `root3`, 40 runs	+0.808	$[+0.411, +1.205]$	see §4

Why the correlation is called fragile. It rests on one point.

Subset	Pearson r	95% CI	Spearman
All 9 maps	+0.783	$[+0.247, +0.952]$	+0.700
Dropping `root5` (whose own CI spans zero)	+0.558	$[-0.242, +0.906]$	+0.571
Polynomials only	+0.766	$[+0.030, +0.963]$	+0.857

Removing a single map moves the interval across zero. And $\bar{n}$ is itself correlated with degree ($r = +0.48$ across the seven polynomials), so with nine maps the two cannot be separated. Most tellingly: **among the five maps with $\bar{n} = 1.00$ exactly, the gain still spans a factor of** 7.3 (0.0050 to 0.0367). Injectivity held perfectly fixed, something else moves the gain sevenfold.

What the null results could and could not have seen. The degree-3 interval has half-width 0.0090 against a gain of about 0.020, so it would only have detected a bijectivity effect larger than roughly **45% of the gain itself**. The observed 17% difference is well inside the noise. Resolving it needs the standard error down by 2.7×, which is **∼240 fits per map** instead of 32 — about eight times the compute, and still only two matched pairs.

The honest summary. We can say with confidence that the gain is real everywhere, that it is smallest for the linear map, that it survives on bijective maps, and that it does not wash out with data. We *cannot* currently say whether bijectivity has a modest secondary effect: the controlled comparison is null but underpowered, and the across-map correlation is positive but fragile and confounded with degree.

One sentence in the paper should be adjusted. "Even a nonlinear composition adds no information advantage when it is invertible" is correct about *information*, but it reads as a prediction that invertible maps will show no gain. They do show one. On this evidence, how nonlinear g is predicts the gain better than whether it is invertible.

10 What would settle it

Section efsec:channels settles it for every g whose fibers are finite: Model 3 is the third arm, and enumerating plausible histories computes it exactly, so the two channels are separated without a sampler and without replicates. Three things remain open.

eginenumerate

extbfMaps with infinite fibers. extttrelu, the clipped pair and the two-dimensional analogues cannot be enumerated, and for those Model 3 has to be sampled rather than computed.

extbfWhether the split survives optimisation. The decomposition is of posteriors at a fixed design. Whether a 3:1 ratio between the channels translates into the same ratio in final regret after forty adaptive evaluations is a separate question.

extbfSeveral intermediates. Everything here is $k = 1$. With $k \geq 2$ a map can destroy information without folding, which is the only way to vary the information channel while holding model difficulty fixed *in a fitted comparison* rather than an exact one.

11 Does this carry over to multiple objectives?

Everything above is single-objective, while the benchmark table the whole question comes from is measured in hypervolume. So the last experiment puts the same contrast into a two-objective problem and uses the paper's own machinery — independent GPs, `qLogEHVI`, the same solvers — rather than anything built for this note.

The structure. One objective carries the map under test; the other is held fixed.

$$x \longrightarrow \begin{cases} h_0 \longrightarrow g_1 \longrightarrow f_1 & g_1 \text{ varies} \\ h_1, h_2, h_3, h_4 \longrightarrow g_2 \longrightarrow f_2 & g_2 \text{ fixed (Ackley)} \end{cases}$$

The composite model observes all five intermediates; the direct model sees only (f_1, f_2). The input dimension, the five inner functions, g_2 and the initial designs are identical across conditions, so any difference between rows is attributable to the structure of g_1 alone.

Two details matter for the comparison to mean anything. Each objective is normalised on a fixed reference set, because hypervolume is not scale-free and z^2 and z^3 otherwise span very different ranges. And all conditions share *one* reference point — taken as the element-wise worst across them — since letting each pick its own would make the hypervolumes incommensurable, which is exactly the comparison being made.

Conditions. The same three regimes as the single-objective study:

g_1		what changes
z	linear	the direct model is *correct*; nothing is lost
z^3, $\sinh z$	bijective, nonlinear	wrong model shape, nothing lost
z^2	not bijective	wrong model shape *and* the sign of h_0 is lost

So linear → bijective isolates misspecification, and bijective → square adds information loss at comparable nonlinearity.

Result. Four function draws × eight seeds, 40 adaptive evaluations, 32 paired runs per condition.

g_1		HV direct	HV composite	Gain (%)	median
z	linear	4.338	5.046	$+17.3 \pm 2.1$	+15.6
$\sinh z$	bijective	5.361	6.427	$+24.2 \pm 5.7$	+23.4
z^3	bijective	5.033	6.267	$+30.8 \pm 6.4$	+22.6
z^2	not bijective	3.935	4.905	$+25.8 \pm 2.2$	+24.8

Composite wins in every condition and on 26 to 31 of 32 runs, which is the first thing worth saying: the advantage is large and robust in hypervolume, as in the benchmark table.

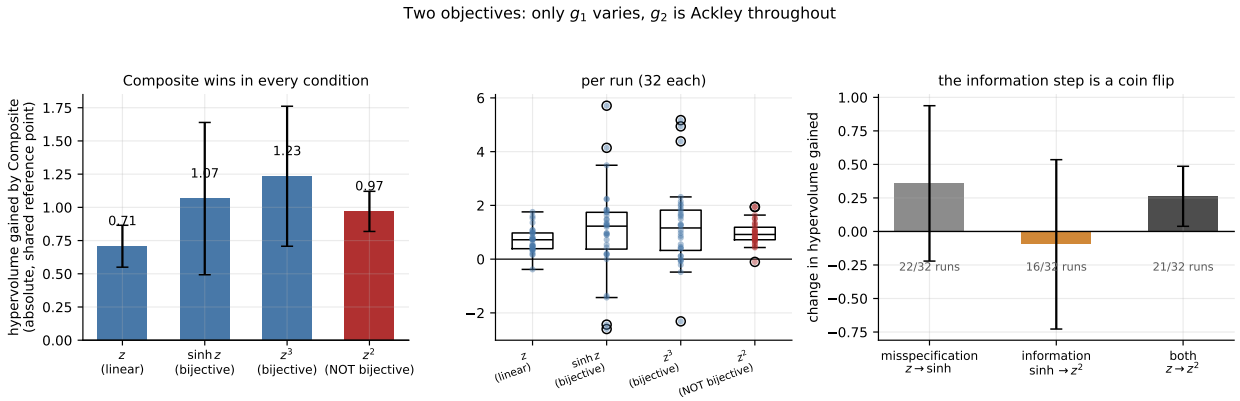


Figure 7: Absolute hypervolume gained by Composite, not percentage. The reference point is shared across conditions but the direct arm's baseline is not (3.94 to 5.36), so a percentage flatters whichever condition the direct arm does worst on — which happens to be the non-bijective one, the very row under test. Left: Composite wins everywhere. Middle: the per-run spread, which is what makes the contrasts hard. Right: the misspecification step is positive on 22/32 runs, the information step an exact coin flip.

But the channels do not separate here. Pairing by draw and seed:

Contrast	Mean	Median	Sign test
with sinh *as the bijective partner*			
$z \to \sinh$ (misspecification only)	$+6.85 \pm 5.97$	$+7.43$	$21/32$, $p = 0.11$
$\sinh \to z^2$ (adds information)	$+1.63 \pm 6.63$	$+5.18$	$20/32$, $p = 0.22$
with z^3 *as the bijective partner*			
$z \to z^3$	$+13.4 \pm 6.6$	-0.7	$16/32$, $p = 1.00$
$z^3 \to z^2$	-5.0 ± 7.3	$+6.5$	$19/32$, $p = 0.38$
$z \to z^2$ (both channels)	$+8.48 \pm 3.22$	$+7.41$	$21/32$, $p = 0.11$

With z^3 as the partner the mean and the median disagree and the sign test reads an exact coin flip, because its gain runs from -29% to $+131\%$ and a few seeds carry the mean far above the median. sinh was added as a second bijective partner on the theory that a tamer range would reduce that variance. It did not — its spread, -38% to $+123\%$, is essentially the same — so the variance is a property of hypervolume on these problems rather than of z^3. But it is far less skewed, its mean and median agree (24.2 against 23.4), and on that ladder both steps point the same way.

The point estimates match the exact decomposition. On the sinh ladder the two channels come out at $+6.85$ and $+1.63$, which add to the $+8.48$ measured end to end, and sit in a ratio of about 4:1. Section 8, by enumerating posteriors rather than running any optimisation, put that ratio at 3 to 7. Two methods with nothing in common agreeing on the split is worth more than either alone — though with standard errors near ± 6 on each component, the agreement is of point estimates, not a confirmation.

Absolute or relative? Reporting the gain as a percentage, as the benchmark table does, turns out to matter here. The reference point is shared, so absolute hypervolume differences are comparable across conditions — but the direct arm's baseline is not, running from 3.94 on z^2 to 5.36 on $\sinh z$. An identical absolute gain therefore becomes a larger percentage on z^2, the one non-bijective condition. On the comparable measure the information step reverses sign:

Contrast	Percentage	Absolute HV
$z \to \sinh$ (misspecification)	$+6.85 \pm 5.97$, $21/32$	$+0.358 \pm 0.296$, $22/32$, $p = 0.05$
$\sinh \to z^2$ (information)	$+1.63 \pm 6.63$, $20/32$	-0.096 ± 0.322, $16/32$, $p = 1.00$

On absolute hypervolume the misspecification step is borderline significant and the information step is an exact coin flip — which is more consistent with the exact decomposition, not less. The same caution applies to reading percentage gains across benchmarks whose hypervolume baselines differ.

Why hypervolume is a blunter instrument. The single-objective study resolved these channels because CRPS is a per-point score averaged over 512 held-out points. Hypervolume is one number per run, produced by a stochastic forty-step optimisation, with a heavy right tail. The standard errors here are roughly five times worse relative to the effect. Resolving the bijective-to-square step at this noise level needs about $13\times$ more runs, roughly 400 per condition.

That is a statement about the *measurement*, not about the effect. The exact decomposition of Section 8 does not depend on any of this: it is computed, not estimated, and it already shows the

information channel to be the smaller of the two. The multi-objective experiment is consistent with that and simply lacks the power to confirm it independently.

What this adds anyway. Two things that do not need the contrast to separate. First, the large composite advantage survives in hypervolume on a problem where the inner functions are drawn from the surrogate's own prior, so it is not an artefact of any particular benchmark. Second, the advantage is substantial *even when* g_1 *is linear* (+17.3%, 31/32 runs) — a case where the proposition guarantees that observing h conveys no information at all. Whatever is producing that gain, it is not information.

12 The functions, written out

Every map in the sweep is $g_s(z) = \prod_{i=1}^{m}(z - s\,a_i)$ with $a = \text{linspace}(-2, 2, m)$. The code `spr<m>s<100s>` records the degree m and the spread s; at $s = 0$ every root collapses to the origin and the map is simply z^m. The last column is the measured average number of h values consistent with one observation.

Code	Description	Roots	$g(z)$	$\bar{n}$
`spr3s000`	degree 3, spread 0	all at 0	z^3	**1.00**
`spr3s025`	degree 3, spread 0.25	$-0.50, +0.00, +0.50$	$z^3 - 0.25z$	1.88
`spr3s050`	degree 3, spread 0.5	$-1.00, +0.00, +1.00$	$z^3 - z$	2.51
`spr3s075`	degree 3, spread 0.75	$-1.50, +0.00, +1.50$	$z^3 - 2.25z$	2.82
`spr3s100`	degree 3, spread 1	$-2.00, +0.00, +2.00$	$z^3 - 4z$	2.96
`spr4s000`	degree 4, spread 0	all at 0	z^4	2.00
`spr4s025`	degree 4, spread 0.25	$-0.50, -0.17, +0.17, +0.50$	$z^4 - 0.278z^2 + 0.00694$	2.81
`spr4s050`	degree 4, spread 0.5	$-1.00, -0.33, +0.33, +1.00$	$z^4 - 1.11z^2 + 0.111$	3.43
`spr4s075`	degree 4, spread 0.75	$-1.50, -0.50, +0.50, +1.50$	$z^4 - 2.5z^2 + 0.562$	3.75
`spr4s100`	degree 4, spread 1	$-2.00, -0.67, +0.67, +2.00$	$z^4 - 4.44z^2 + 1.78$	3.93
`spr5s000`	degree 5, spread 0	all at 0	z^5	**1.00**
`spr5s025`	degree 5, spread 0.25	$-0.50, -0.25, +0.00, +0.25, +0.50$	$z^5 - 0.312z^3 + 0.0156z$	2.33
`spr5s050`	degree 5, spread 0.5	$-1.00, -0.50, +0.00, +0.50, +1.00$	$z^5 - 1.25z^3 + 0.25z$	3.39
`spr5s075`	degree 5, spread 0.75	$-1.50, -0.75, +0.00, +0.75, +1.50$	$z^5 - 2.81z^3 + 1.27z$	4.06
`spr5s100`	degree 5, spread 1	$-2.00, -1.00, +0.00, +1.00, +2.00$	$z^5 - 5z^3 + 4z$	4.50
`spr6s000`	degree 6, spread 0	all at 0	z^6	2.00
`spr6s025`	degree 6, spread 0.25	$-0.50, -0.30, -0.10, +0.10, +0.30, +0.50$	$z^6 - 0.35z^4 + 0.0259z^2 - 0.000225$	3.32
`spr6s050`	degree 6, spread 0.5	$-1.00, -0.60, -0.20, +0.20, +0.60, +1.00$	$z^6 - 1.4z^4 + 0.414z^2 - 0.0144$	4.38
`spr6s075`	degree 6, spread 0.75	$-1.50, -0.90, -0.30, +0.30, +0.90, +1.50$	$z^6 - 3.15z^4 + 2.1z^2 - 0.164$	5.10
`spr6s100`	degree 6, spread 1	$-2.00, -1.20, -0.40, +0.40, +1.20, +2.00$	$z^6 - 5.6z^4 + 6.63z^2 - 0.922$	5.54

Bold $\bar{n}$ marks the two injective members, z^3 and z^5. Even powers fold at the origin, so z^4 and z^6 are two-to-one and are not injective however the roots are placed.

Caveats

- One function family, $d = 6$, single objective, noiseless. The paper's hypervolume benchmarks are none of those things.
- `root5`'s error bar is enormous (± 0.068 on $+0.113$); a few outlier fits dominate it. The degree-3 pair is the trustworthy comparison.

- $\bar{n}$ counts branches without weighting them by how plausible they are, so a map with a technically reachable but very unlikely second branch looks more ambiguous than it behaves.