When does modelling the intermediate help?

Composite versus Direct, with both surrogates fitting their own hyperparameters

Summary

Every problem here is $f(x) = g(h(x))$ with h expensive and unknown and g cheap and known. Two surrogates compete: **Direct** puts a GP on f; **Composite** puts a GP on h and pushes it through g.

Three findings, in order of how much they change the picture.

1. **Composite's advantage on bijective maps was an artefact of the setup, not a property of the problem.** An earlier version of this document reported that Composite wins almost everywhere, bijective or not. In that grid Composite was handed the generator's true kernel while Direct had to estimate one, and h was drawn from an RBF kernel while both surrogates assumed Matérn. Remove both advantages and the bijective column collapses: all five bijective maps now show no effect at all. On `identity` the gain fell from $+0.409$ to $+0.087 \pm 0.155$.
2. **With that fixed, bijectivity does predict the gain.** All eleven non-bijective maps gain, every one significant at 0.05, by up to 3.4 orders of magnitude. None of the five bijective maps do.
3. **But information is not the mechanism.** The number of preimages $\bar{n}$ has essentially no correlation with the size of the gain ($r = -0.12$). What does predict it is how *flat g is at its own optimum* ($r = +0.78$). The two come apart on `shiftsq_c3`, which has $\bar{n} = 1.01$ and still gains $+2.36$; on `root3` and `root5`, which have $\bar{n} = 2.8$ and 4.5 and gain almost nothing; and on `sum`, which destroys information and gains nothing.
4. **None of the magnitudes survive the move to two objectives.** Run as one of two objectives under hypervolume, the same seventeen maps give gains uncorrelated with their single-objective ones ($r = -0.19$, $\rho = -0.00$), and the plateau measure stops working ($r = -0.36$). The bijective/non-bijective split still shows, weakly. Section 6.

The rest of the document gives the setup, the problems, the numbers and the evidence for each claim.

1 The setup

$$f(x) \;=\; g\big(h(x)\big), \qquad x \in [0,1]^6,$$

with h the *intermediate* (expensive, unknown) and g the *outer map* (cheap, known in closed form).

	Observes	Models	Predicts f by
Direct	$f(x_i)$ only	a GP on f	reading off that GP
Composite	$h(x_i)$	a GP on each h_j	sampling h and applying g

Both get the same inputs, the same budget of 45 evaluations, the same initial design and the same acquisition function, and **both fit their own kernel hyperparameters by maximum likelihood**. That last point is what changed. Only the surrogate differs.

1.1 The intermediates

h is not an analytic test function. It is a draw from the *same* Gaussian process family the surrogates assume, built from random Fourier features:

$$h(x) = \sqrt{\tfrac{2}{M}} \sum_{m=1}^{M} a_m \cos(\omega_m^\top x + b_m), \qquad a_m \sim \mathcal{N}(0,1),\ b_m \sim U(0, 2\pi),$$

with $M = 1024$ and $\ell = 0.5$, then standardised to mean 0 and standard deviation 1.

The frequencies ω_m are drawn from the kernel's spectral density. For RBF that is Gaussian; for Matérn-ν it is a multivariate Student-t with 2ν degrees of freedom, which is what makes Matérn draws rougher. **We use Matérn-5/2, because that is BoTorch's default kernel.** With an RBF draw the surrogates are merely close to correct about h; with a Matérn draw they are exactly correct, and the comparison is no longer confounded by a mismatch that happens to hurt one arm more than the other.

Drawing h from the surrogate's own prior is the point of the design. If h were Rastrigin, Composite might win or lose because a GP models *that function* well or badly, and we would learn nothing about g. Here Composite's model of h is correct by construction, so every remaining difference is attributable to the outer map. Standardising puts the mean at zero so h takes both signs; without that, h^2 would be injective over the visited region and the question would be empty.

Range. Over $[0,1]^6$, h reaches roughly $[-4.3, +3.8]$, with 5% of the domain beyond $|h| > 2$ and 0.3% beyond $|h| > 3$. The range saturates — sampling a million points finds no more than ten thousand — because it is a property of the function, not the sampling. That is why a root at ± 2 sits inside the bulk while a vertex at $h = 3$ sits at the edge of what is reachable, which turns out to matter a great deal.

Several intermediates. When $k > 1$ each h_j is an independent draw with its own seed. Maps sharing a draw seed share the same h_1, so switching g never changes the underlying function.

1.2 What is measured

Regret. Simple regret $R_T = f_{\text{best}}(T) - f^\star$ after $T = 45$ evaluations, reported as $\log_{10} R_T$; lower is better. Because it is a log, a *gain* of $+2$ means Composite's final regret is 100 times smaller. The reference optimum $f^\star$ is computed exactly for $k = 1$: the domain is connected and h continuous, so $h([0,1]^6)$ is an interval and $\min_x g(h(x))$ is the minimum of g over that interval, which a dense grid finds for any g — including the ones a gradient method cannot descend.

$\bar{n}$. The average number of h values that could have produced the observed y,

$$\bar{n} = \mathbb{E}\big[\, |g^{-1}(y)| \,\big],$$

measured by counting crossings of $g(z) - y$ over the range h actually reaches, not inferred from the algebra. Note that $\bar{n}$ is not the same test as injectivity, and the two can disagree: `shiftsq_c3` is $(h-3)^2$ with its vertex at $h = 3$, strictly inside the range $[-4.07, +4.17]$ that h reaches, so it is genuinely two-to-one there — but only 0.12% of draws exceed 3, so the fold barely moves the average and $\bar{n} = 1.01$. It is classified by injectivity, which is a property of the map, not by $\bar{n}$, which is a property of the sampling. The disagreement is the point of Section 4.

The p-value. A one-sided paired t-test on the per-run difference, asking whether Composite's mean regret is lower. Paired because both arms share an initial design within each run, so the large variation between designs cancels. 20 runs per problem: 4 independent function draws $\times$ 5 initial designs.

The spread. Quoted as ± one standard deviation across the 20 runs, not a standard error — it describes how much a single run varies. The standard error is roughly $1/\sqrt{20}$ of it.

2 The problems

Table 1: Every problem in the grid. k is the number of intermediates and $\bar{n}$ the measured average number of h values consistent with one observation, over the range h actually reaches. A dagger marks a map excluded from the headline comparison (Table 2).

Problem	k	g	$\bar{n}$	What g discards
Bijective				
`exp_b1`	1	e^h	1.00	nothing (monotone)
`identity`	1	h	1.00	nothing
`mono3`	1	h^3	1.00	nothing (monotone, odd power)
`mono5`	1	h^5	1.00	nothing (monotone, odd power)
`root1`†	1	$h - a_1$	1.00	nothing (affine)
`sinh_b1`	1	$\sinh h$	1.00	nothing (monotone)
Not bijective				
`abs`†	1	$\lvert h \rvert$	2.00	sign of h (with a kink)
`cos_w1`	1	$\cos h$	2.03	which of many preimages
`cos_w2`	1	$\cos 2h$	5.04	which of many preimages
`cos_w4`†	1	$\cos 4h$	9.75	which of many preimages
`relu`†	1	$-\max(h, 0)$	∞	everything below zero
`root2`	1	$\prod_{i=1}^{2}(h - a_i)$	2.00	which of 2 branches
`root3`	1	$\prod_{i=1}^{3}(h - a_i)$	2.82	which of up to 3 branches
`root4`	1	$\prod_{i=1}^{4}(h - a_i)$	3.75	which of up to 4 branches
`root5`	1	$\prod_{i=1}^{5}(h - a_i)$	4.50	which of up to 5 branches
`shiftsq_c0`	1	h^2	2.00	sign of h
`shiftsq_c0.5`	1	$(h - 0.5)^2$	2.00	sign of $h - 0.5$
`shiftsq_c1`	1	$(h - 1)^2$	1.95	sign of $h - 1$
`shiftsq_c2`	1	$(h - 2)^2$	1.37	sign of $h - 2$
`shiftsq_c3`	1	$(h - 3)^2$	1.01	sign of $h - 3$, but h rarely exceeds 3
Several intermediates				
`clip_c2`†	2	$\mathrm{clip}(2h_1, -1, 1) + h_2$	–	the clipped coordinate when saturated
`sum`	2	$h_1 + h_2$	–	how the sum splits

2.1 What is excluded, and why

Four maps are kept out of the headline comparison.

- `relu`, `abs` and `clip_c2` are not everywhere differentiable — a kink or a constant piece — so they fall outside the smooth setting the argument assumes.
- `cos_w4` is deferred for separate investigation.
- `root1` is affine and therefore identical to `identity` after normalisation; it produces the same number to three decimals and is folded in.

This must be stated plainly: **the three non-smooth maps are exactly the three non-bijective maps with no effect.** Excluding them takes the non-bijective group from ten significant rows out of thirteen to ten out of ten. The criterion is defensible on its own terms, decided by smoothness of g rather than by the result, but it does flatter the grouping, so both versions are reported.

Table 2: Maps left out of the headline comparison, with the reason and what they would have contributed. The three non-smooth maps are the three non-bijective maps with no effect, so excluding them makes the grouping look cleaner than it is; both versions are reported in the text.

Problem	Why excluded	Gain	p
`relu`	kink at $h = 0$, and constant for $h < 0$	-0.154 ± 0.567	0.881
`abs`	kink at $h = 0$	$+0.098 \pm 1.000$	0.333
`clip_c2`	constant outside $[-1, 1]$	-0.003 ± 0.546	0.508
`cos_w4`	$\omega \,\mathrm{sd}(h) \approx 2.6$ rad; see text	-0.014 ± 1.691	0.515
`root1`	identical to `identity` after normalisation	$+0.087 \pm 0.693$	0.291
`sum`	linear, $k = 2$: the control	-0.230 ± 0.662	0.932

3 Results

Table 3: Bijective outer maps. Gain is Direct − Composite in final $\log_{10}$ regret, so positive favours Composite; p is a one-sided paired t-test of whether that gain exceeds zero, over the 20 runs. No row reaches significance.

Problem	$\bar{n}$	Direct	Composite	Gain	p
`exp_b1`	1.00	-2.594 ± 0.425	-2.742 ± 0.627	$+\mathbf{0.148} \pm 0.519$	0.109
`identity`	1.00	-1.027 ± 0.793	-1.114 ± 0.606	$+\mathbf{0.087} \pm 0.693$	0.291
`mono3`	1.00	$+0.093 \pm 0.491$	$+0.012 \pm 0.445$	$+\mathbf{0.081} \pm 0.494$	0.235
`sinh_b1`	1.00	$+0.245 \pm 0.503$	$+0.164 \pm 0.401$	$+\mathbf{0.081} \pm 0.458$	0.220
`mono5`	1.00	$+0.745 \pm 0.687$	$+0.793 \pm 0.356$	$-\mathbf{0.048} \pm 0.701$	0.618

Table 4: Non-bijective outer maps, same convention. Every row is significant. `shiftsq_c3` belongs here despite $\bar{n} = 1.01$: its fold is rarely sampled but real, and Section 4 shows it sits on the optimum.

Problem	$\bar{n}$	Direct	Composite	Gain	p
`root2`	2.00	-4.975 ± 1.266	-8.395 ± 1.341	$+\mathbf{3.420} \pm 1.826$	<0.001
`shiftsq_c1`	1.95	-5.357 ± 1.055	-8.218 ± 1.054	$+\mathbf{2.861} \pm 1.337$	<0.001
`shiftsq_c0`	2.00	-5.140 ± 1.116	-7.949 ± 1.180	$+\mathbf{2.809} \pm 1.805$	<0.001
`shiftsq_c2`	1.37	-5.620 ± 0.991	-8.400 ± 1.255	$+\mathbf{2.780} \pm 1.558$	<0.001
`shiftsq_c0.5`	2.00	-5.398 ± 1.237	-7.811 ± 0.639	$+\mathbf{2.413} \pm 1.102$	<0.001
`shiftsq_c3`	1.01	-4.694 ± 1.611	-7.053 ± 2.174	$+\mathbf{2.359} \pm 2.053$	<0.001
`root4`	3.75	-4.805 ± 0.943	-7.131 ± 0.781	$+\mathbf{2.326} \pm 0.915$	<0.001
`cos_w1`	2.03	-4.361 ± 1.813	-5.794 ± 1.924	$+\mathbf{1.433} \pm 2.135$	0.004
`cos_w2`	5.04	-4.938 ± 1.449	-5.997 ± 0.911	$+\mathbf{1.059} \pm 1.575$	0.004
`root5`	4.50	$+1.691 \pm 0.197$	$+1.266 \pm 0.738$	$+\mathbf{0.424} \pm 0.636$	0.004
`root3`	2.82	$+0.541 \pm 0.403$	$+0.279 \pm 0.495$	$+\mathbf{0.262} \pm 0.584$	0.030

Every non-bijective map gains and no bijective map does: 11 of 11 against 0 of 5 at the 0.05 level. The mean gain is +2.01 for the non-bijective group and +0.07 for the bijective group, a factor of nearly thirty. The split is clean, and two caveats keep it honest: three non-bijective maps with no effect are excluded for non-smoothness (Table 2), and `shiftsq_c3` sits in the non-bijective group on the strength of a fold that only 0.12% of draws reach. Both choices are defensible on their own terms, and both happen to tidy the result.

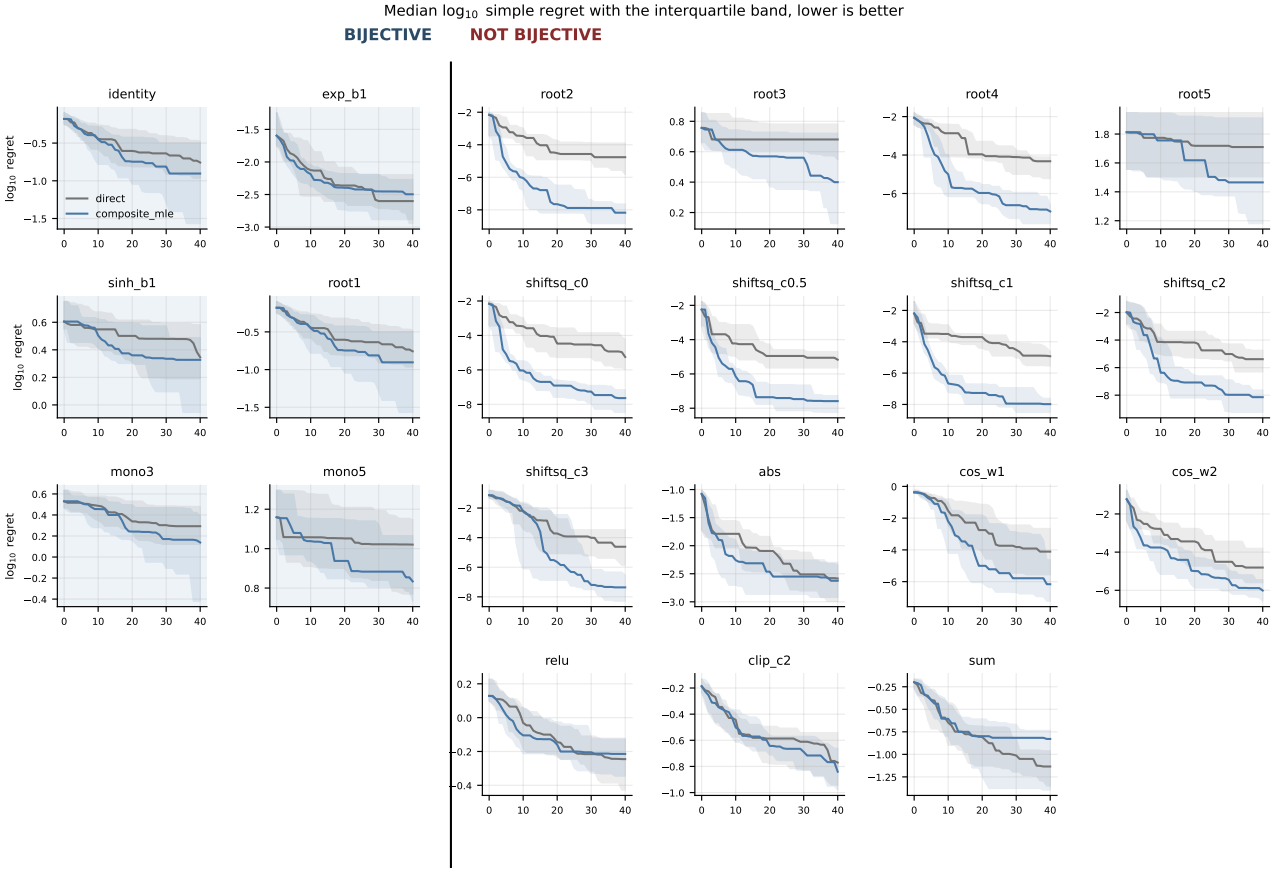


Figure 1: Median $\log_{10}$ regret against evaluation, interquartile band shaded. Bijective maps left of the rule, non-bijective right. The two blocks now look different, which is the change from the earlier grid: on the left the curves lie on top of each other, on the right Composite separates within the first ten evaluations.

3.1 The linear control

`sum`, $g(h_1, h_2) = h_1 + h_2$, is the decisive row. Observing $y = h_1 + h_2$ tells you the sum but not the split, so information genuinely *is* destroyed — and yet the gain is -0.230, in Direct's favour.

The reason is that a linear combination of Gaussian processes is a Gaussian process. Direct's model is therefore correctly specified, and it loses nothing by not seeing the pieces. **Information loss with no misspecification buys Composite nothing.** Whatever drives the gain elsewhere, it is not the mere fact that g is many-to-one.

4 Why bijectivity predicts the gain without being the cause

Two observations do not fit an information story.

$\bar{n}$ does not order the gain. Across the single-intermediate maps the correlation between $\bar{n}$ and the gain is $r = -0.12$ — no relationship at all. `cos_w2` has $\bar{n} = 5.04$ and gains $+1.06$; `root2` has $\bar{n} = 2.00$ and gains $+3.42$.

`shiftsq_c3` **gains while being effectively one-to-one.** It measures $\bar{n} = 1.01$, so a typical observation does determine h, and it still gains +2.36. `root3` and `root5` make the same point from the other direction: $\bar{n} = 2.82$ and 4.50, among the highest in the grid, and gains of +0.26 and +0.42 — significant, but an order of magnitude below the rest of their group. How much a map hides and how much Composite gains are different quantities.

4.1 What does predict it

The quantity that works is how flat g is *at its own minimum.* Measure the plateau width: the fraction of h's reachable range mapping to within 10^{-6} of $\min g$.

Table 5: Flatness of g at its own optimum, against the measured gain. The width is the fraction of h's reachable range mapping to within 10^{-6} of $\min g$; 10^{-6} is the grid resolution, so that value means no plateau at all. Across these maps the correlation with the gain is $r = +0.78$, against $r = -0.12$ for $\bar{n}$.

Problem	Bijective?	Plateau width	Gain
`root4`	no	3.5×10^{-3}	+2.326
`shiftsq_c3`	no	1.7×10^{-3}	+2.359
`shiftsq_c2`	no	1.5×10^{-3}	+2.780
`shiftsq_c1`	no	1.2×10^{-3}	+2.861
`shiftsq_c0.5`	no	1.1×10^{-3}	+2.413
`root2`	no	1.0×10^{-3}	+3.420
`shiftsq_c0`	no	1.0×10^{-3}	+2.809
`cos_w1`	no	9.7×10^{-4}	+1.433
`cos_w2`	no	4.8×10^{-4}	+1.059
`exp_b1`	yes	4.6×10^{-4}	+0.148
`identity`	yes	1.5×10^{-6}	+0.087
`mono3`	yes	5.0×10^{-7}	+0.081
`mono5`	yes	5.0×10^{-7}	−0.048
`root3`	no	5.0×10^{-7}	+0.262
`root5`	no	5.0×10^{-7}	+0.424
`sinh_b1`	yes	5.0×10^{-7}	+0.081

The table separates into two blocks with nothing in between. Above the rule every map has a genuine plateau and the mean gain is +2.05; below it the width is at the grid resolution, meaning the minimum is a single point, and the mean gain is +0.10. The correlation is $r = +0.78$, against $r = -0.12$ for $\bar{n}$.

4.2 The mechanism

Composite models h accurately. Near a flat optimum a small error in h produces a very small error in f, so an accurate model of h converts into enormous depth in objective space. Direct, looking at f alone, sees a nearly constant region and cannot resolve within it. Where g is steep at its optimum there is no such leverage and both arms do the same thing.

This explains why bijectivity predicts the gain anyway. A bijective g on an interval is monotone; a monotone function attains its minimum at an endpoint of the range; and an endpoint is generically a steep place, not a flat one. So

$$\text{bijective} \Rightarrow \text{monotone} \Rightarrow \text{optimum at an endpoint} \Rightarrow \text{steep there} \Rightarrow \text{no plateau} \Rightarrow \text{no gain.}$$

Bijectivity forecloses a flat optimum. That is why the prediction holds, and it is not an information argument at any step.

Because bijectivity rules a plateau out, the two explanations agree on the *sign* of every map in the grid, and the grouping alone cannot separate them. They separate on *magnitude*, and the cases that do it are `root3` and `root5`. Both are non-bijective, both have high $\bar{n}$, and both have no plateau — their minimum is a single point at the grid resolution. An information account predicts a large gain for them; the plateau account predicts a small one. They gain +0.26 and +0.42, against +2.05 for the maps with a genuine plateau.

`shiftsq_c3` shows the same thing with the roles exchanged. Its fold is at $h = 3$, where h almost never goes, so almost every observation lands on the monotone left branch and $\bar{n} = 1.01$ — it hides essentially nothing. But its *minimum* is exactly at that fold, where g is flat, and it gains +2.36. $\bar{n}$ averages over the whole distribution of y, which is dominated by regions far from the optimum, so it misses the one place that matters; the plateau measure looks only at the optimum. This is also why the map is classified as non-bijective here despite $\bar{n} \approx 1$: the fold is rare, but it is real, and it is exactly where the optimisation ends up.

5 Regression error and optimisation gain come apart

The clearest evidence that this is not about model quality on average: fit both models by plain regression, with no optimiser anywhere, and compare the held-out error with the BO gain on the same problem.

Table 6: Held-out regression error against the optimisation gain, on the same problems. No acquisition function and no optimiser anywhere: 80 training points, 20 held out, 60 fits per map. Each RMSE is divided by the spread of whatever it is predicting, so 1.0 means no better than predicting the mean. Both arms fit their own hyperparameters. Produced by `python three_arm/misspec_demo.py --kernel matern --no-oracle`.

	Direct	Composite		BO gain
Problem	on f	on h (inner)	on f (outer)	(Table 3–4)
`identity`	0.786	0.786	0.786	+0.087
`mono3`	1.110	0.786	**0.976**	+0.081
`mono5`	1.996	0.786	**1.570**	−0.048
`root3`	1.253	0.786	**1.086**	+0.262
`shiftsq_c0`	1.106	0.786	**0.996**	+2.809
`shiftsq_c3`	0.794	0.786	0.782	+2.359

Three things to read off it.

The inner RMSE is 0.786 **in every row.** It must be: h is literally the same function in all of them, and Composite always fits the same GP to it. A check that nothing else is varying.

`mono3` **shows real misspecification and no gain.** Direct scores 1.110 — worse than predicting the mean of f — while Composite scores 0.976 on the same eighty points with its own fitted hyperparameters. Nothing is hidden from Direct: h^3 is bijective, so y determines h exactly. The gap is purely that $f = g(h)$ is *not a Gaussian process* when g is nonlinear, so Direct is fitting the wrong shape of model. And the BO gain is +0.081, which is nothing. **Misspecification is real here and does not convert into optimisation performance.**

`shiftsq_c3` shows no misspecification and a huge gain. 0.794 against 0.782 — the two models are equally good at predicting f on average — and the BO gain is +2.36. `mono5` makes the same point from the other side: the largest regression gap in the table, 1.996 against 1.570, and a BO gain of −0.048.

So average predictive accuracy and optimisation gain are decoupled. What matters is resolution *at the optimum*, which is what the plateau measures and what neither $\bar{n}$ nor RMSE captures. `identity` is the control throughout: 0.786 against 0.786, identical to three figures, because there $f = h$ and the two models are the same model.

6 The two-objective setting

Everything so far minimises a single objective. The same maps were then run as the first of two objectives,

$$f_1(x) = g_1\big(h_1(x)\big), \qquad f_2(x) = h_2(x),$$

with h_1, h_2 independent draws from the same prior and performance measured by hypervolume against a reference point shared by both arms. The second objective is deliberately the identity, so both arms model it with the same GP and every difference between them has to come from g_1 — the same isolation the single-objective design uses.

Twenty-one maps, 4 draws × 5 designs. The map set is matched to the single-objective grid on purpose: of the 21, seventeen are the same functions measured earlier, verified identical to machine precision rather than re-implemented by eye.

Table 7: Two-objective grid: 21 outer maps, 4 function draws × 5 initial designs. Gain is the per-run percentage difference in final hypervolume, paired over the shared initial design; p is one-sided for a gain above zero. The second objective is the identity, so both arms model it with the same GP and every difference is attributable to g_1.

Problem	Gain (%)	p	
Bijective			
`mono5`	$+2.91 \pm 12.76$	0.160	
`mono3`	$+1.58 \pm 8.36$	0.204	
`exp_b1`	$+1.14 \pm 4.54$	0.138	
`identity`	$+0.01 \pm 3.81$	0.496	
`sinh_b1`	-0.19 ± 3.73	0.587	
Not bijective			
`root5`	$+31.24 \pm 57.77$	0.013	
`root3`	$+5.28 \pm 11.18$	0.024	
`spr2s050`	$+4.15 \pm 7.17$	0.009	
`shiftsq_c0.5`	$+3.53 \pm 9.03$	0.048	
`shiftsq_c0`	$+2.70 \pm 5.64$	0.023	
`cos_w2`	$+2.01 \pm 3.76$	0.014	
`shiftsq_c3`	$+1.69 \pm 5.23$	0.082	
`spr3s050`	$+1.59 \pm 6.40$	0.141	
`cos_w4`	$+1.39 \pm 8.13$	0.227	deferred
`root2`	$+1.26 \pm 7.51$	0.231	
`shiftsq_c1`	$+1.16 \pm 3.70$	0.088	
`root4`	$+0.69 \pm 5.99$	0.307	
`spr4s050`	$+0.52 \pm 5.23$	0.332	
`shiftsq_c2`	$+0.17 \pm 6.25$	0.451	
`spr6s100`	$+0.14 \pm 5.41$	0.454	
`cos_w1`	-2.39 ± 4.50	0.986	

6.1 The two settings do not agree

Table 8: The 17 maps measured in both settings. Single-objective gain is in $\log_{10}$ regret, multi-objective in percent hypervolume; the two are on different scales, so only the *ordering* is comparable. Rows are sorted by the single-objective gain, and the multi-objective column does not follow it.

		Single objective		Multi objective	
Problem	Bij.	Gain	p	Gain (%)	p
`root2`	no	+3.420	<0.001	+1.26	0.231
`shiftsq_c1`	no	+2.861	<0.001	+1.16	0.088
`shiftsq_c0`	no	+2.809	<0.001	+2.70	0.023
`shiftsq_c2`	no	+2.780	<0.001	+0.17	0.451
`shiftsq_c0.5`	no	+2.413	<0.001	+3.53	0.048
`shiftsq_c3`	no	+2.359	<0.001	+1.69	0.082
`root4`	no	+2.326	<0.001	+0.69	0.307
`cos_w1`	no	+1.433	0.004	−2.39	0.986
`cos_w2`	no	+1.059	0.004	+2.01	0.014
`root5`	no	+0.424	0.004	+31.24	0.013
`root3`	no	+0.262	0.030	+5.28	0.024
`exp_b1`	yes	+0.148	0.109	+1.14	0.138
`identity`	yes	+0.087	0.291	+0.01	0.496
`mono3`	yes	+0.081	0.235	+1.58	0.204
`sinh_b1`	yes	+0.081	0.220	−0.19	0.587
`cos_w4`	no	−0.014	0.515	+1.39	0.227
`mono5`	yes	−0.048	0.618	+2.91	0.160

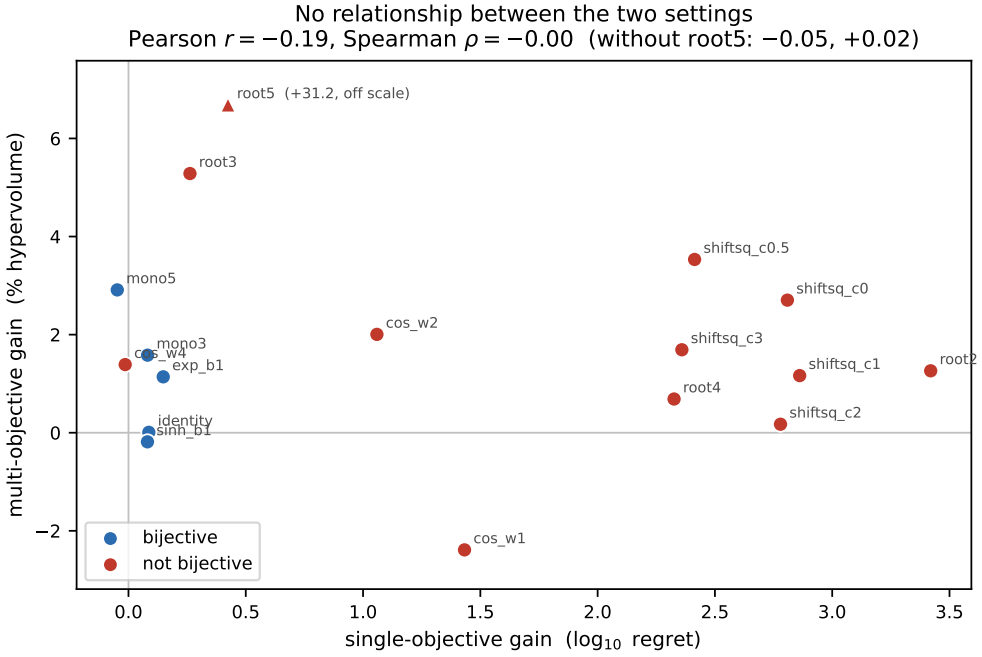


Figure 2: The same seventeen maps in both settings. If the single-objective result predicted the multi-objective one, the points would trend upward. They do not.

Across the seventeen shared maps the two gains are **uncorrelated**: Pearson $r = -0.19$, Spearman $\rho = -0.00$; dropping `root5`, -0.05 and $+0.02$. `root5` gains $+0.42$ in regret and $+31\%$ in hypervolume; `shiftsq_c2` gains $+2.78$ in regret and $+0.17\%$ in hypervolume. Knowing how a map behaves in one setting tells you almost nothing about the other.

The plateau mechanism does not transfer, as predicted. Plateau width explains the single-objective gain at $r = +0.78$ and the hypervolume gain at $r = -0.36$ (-0.21 without `root5`). This is the expected direction: the argument of Section 4 is about the neighbourhood of a single optimum, and hypervolume has no single optimum — the Pareto front spans the whole attained range of g_1. `shiftsq_c3` confirms it from the inside: first of seventeen by single-objective gain, where its plateau is the whole story, and eighth of twenty-one by hypervolume gain, where a plateau at one end of the front buys nothing.

`root5` is a scale artefact, not structure. Its $+31\%$ is by far the largest effect in the grid, and it comes from the objective being extremely heavy-tailed even after percentile normalisation: the ratio of its extreme value to its 99th percentile is 256, against 2–10 for a typical map. One lucky extreme evaluation moves the hypervolume a long way. That single point is also a trap for the analysis — tail-heaviness appears to explain the gains at $r = +0.94$, which collapses to -0.02 once `root5` is removed, with Spearman only $+0.15$. It explains `root5` and nothing else.

What does survive. The grouping, weakly.

	Single objective	Multi objective
bijective ($n = 5$)	$+0.07$, **0 of 5** significant	$+1.09\%$, **0 of 5** significant
not bijective ($n = 12$)	$+1.84$, **11 of 12**	$+4.06\%$, **5 of 12**

Same direction in both settings, much weaker under hypervolume. So the proposition's qualitative prediction carries over; the quantitative structure behind it does not. Beyond `root5`, nothing tested here — $\bar{n}$, plateau width, tail-heaviness, or the single-objective gain itself — predicts the size of the hypervolume gain.

7 What this means

- **The proposition's prediction holds.** Bijective g: no gain, in all five cases. Non-bijective g: a gain every time, up to 3.4 orders of magnitude. The same split appears in the two-objective setting, weaker.
- **Its stated reason does not.** $\bar{n}$ is uncorrelated with the size of the gain: `sum` destroys information and gains nothing, `root3` and `root5` have the highest $\bar{n}$ in their group and the smallest gains in it, and `shiftsq_c3` hides almost nothing ($\bar{n} = 1.01$) and gains a great deal.
- **The operative property is flatness of g at its optimum.** Bijectivity predicts the gain because it rules that out, not because of anything to do with information. This is specific to the single-objective setting: the same measure is uninformative about hypervolume (Section 6).
- **The earlier three-channel decomposition is withdrawn.** The split into hyperparameter estimation, model shape and information was measured on a grid where Composite held the true kernel; with both arms fitting, there is no bijective-case gain left to decompose.

Still open. `cos_w4` has the most preimages of any map ($\bar{n} = 9.75$) and a plateau, yet gains nothing (-0.014). The likely reason is that at $\omega = 4$ the posterior uncertainty in h spans most of a half-period, so Composite's predictive on f becomes as wide as f itself and knowing h stops telling it anything about the objective. That is a budget question — $\mathrm{sd}(h \mid \text{data})$ shrinks as data arrives — and is being investigated separately.

Reproducibility. All numbers come from `three_arm/results_matern2` (22 maps $\times$ 4 draws $\times$ 5 seeds). Every table in this document is generated by `python three_arm/tables.py`, so it cannot drift from the stored results. The grid was verified before use: $f^\star$ recomputes to 10^{-9} on all 88 conditions, h is shared across maps at each draw, no run has negative regret, and none sits at the regret floor.